

# Machine Learning with Graphs

Simona Juvină

National University of Science and Technology POLITEHNICA Bucharest, Romania  
BIOSINF Master Program

May 2025



## 1 Introduction

- Why Graphs?

## 2 Types of Tasks

## 3 Why is Graph Deep Learning Challenging?

## 4 Deep Learning on Graphs

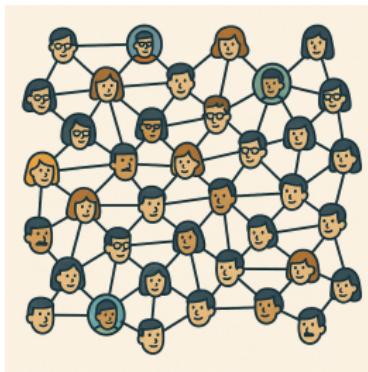
## 5 Classical GNN Layers

# Why Graphs?

- Graphs are a universal language for understanding relationships and interactions.
- What data can we model as graphs?

# Why Graphs?

- Graphs are a universal language for understanding relationships and interactions.
- What data can we model as graphs?



Social Networks



Computer Networks

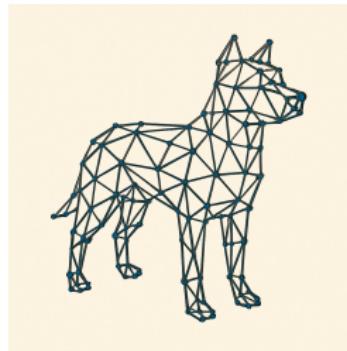


Road Networks

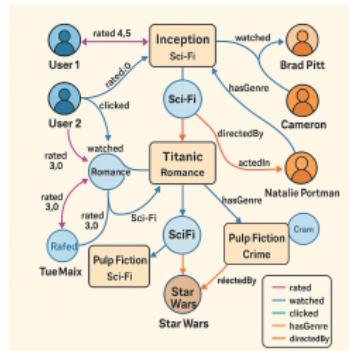
# Why Graphs?



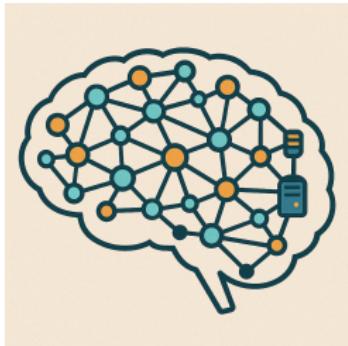
Citation Networks



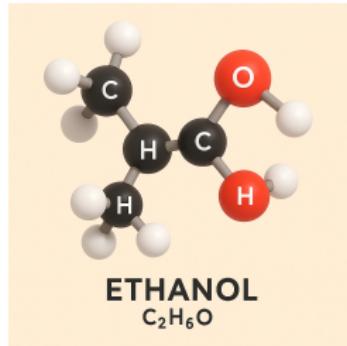
3D Meshes



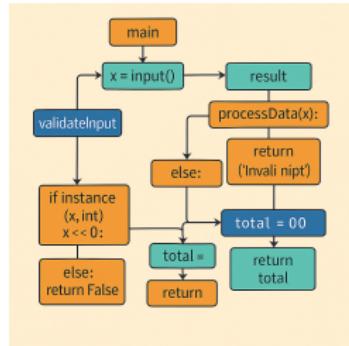
Recommender Graphs



Brain Connectomes



Molecular Graphs

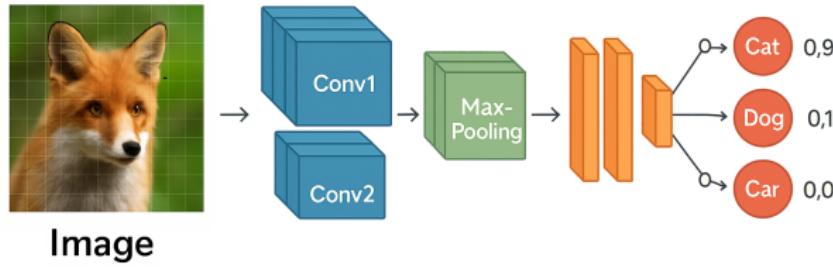


Program Code Graphs

# Why Graphs? Limitations of Traditional Deep Learning

- **Modern deep learning models** — such as **CNNs**, **RNNs**, and **Transformers** — are designed to operate on data organized as **grids** (e.g., images) or **sequences** (e.g., text, time series).
- These formats are:
  - ▶ Well understood
  - ▶ Supported by **highly optimized architectures**
  - ▶ Ideal for domains like computer vision, NLP, and time-series prediction

I cannot believe how long it's taking me to write this course,  
I'm going to go get a beer.



# But There's a Gap...

## Why Traditional Models Fall Short

- Many real-world problems involve **complex relationships** between entities:
  - ▶ Social connections
  - ▶ Biological interactions
  - ▶ Recommender systems
  - ▶ Scientific structures (e.g., molecules)
- These dependencies **don't fit neatly into a grid or sequence.**

# But There's a Gap...

## Why Traditional Models Fall Short

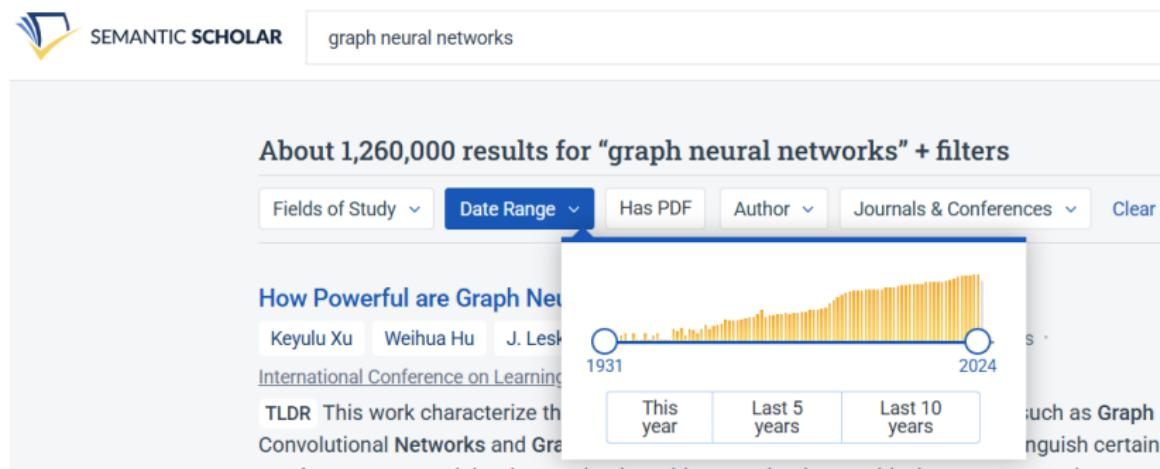
- Many real-world problems involve **complex relationships** between entities:
  - ▶ Social connections
  - ▶ Biological interactions
  - ▶ Recommender systems
  - ▶ Scientific structures (e.g., molecules)
- These dependencies **don't fit neatly into a grid or sequence.**

## Graphs Fill This Gap

***Graphs enable flexible, relational representations where structure is irregular, sparse, and non-Euclidean — and traditional deep learning models struggle.***

# The Rise of Graph Neural Networks

- Research on Graph Neural Networks (GNNs) has grown dramatically over the past decade.
- The figure below shows the increasing number of papers published each year with “Graph Neural Networks” in their title.



# Why GNNs Are More Than Just a Trend

- GNNs combine the structure of graphs with the predictive power of deep learning.
- They can learn sophisticated node and edge representations (embeddings).
- Real-world systems actively use them for tasks requiring relational reasoning.
- Their key strength: capturing both node attributes and relationships between entities.

# Why GNNs Are More Than Just a Trend

- GNNs combine the structure of graphs with the predictive power of deep learning.
- They can learn sophisticated node and edge representations (embeddings).
- Real-world systems actively use them for tasks requiring relational reasoning.
- Their key strength: capturing both node attributes and relationships between entities.

## Is My Problem a Graph Problem?

If your data has natural **connections or interdependencies** between entities — users, products, merchants, documents — it may benefit from a graph-based approach.

## 1 Introduction

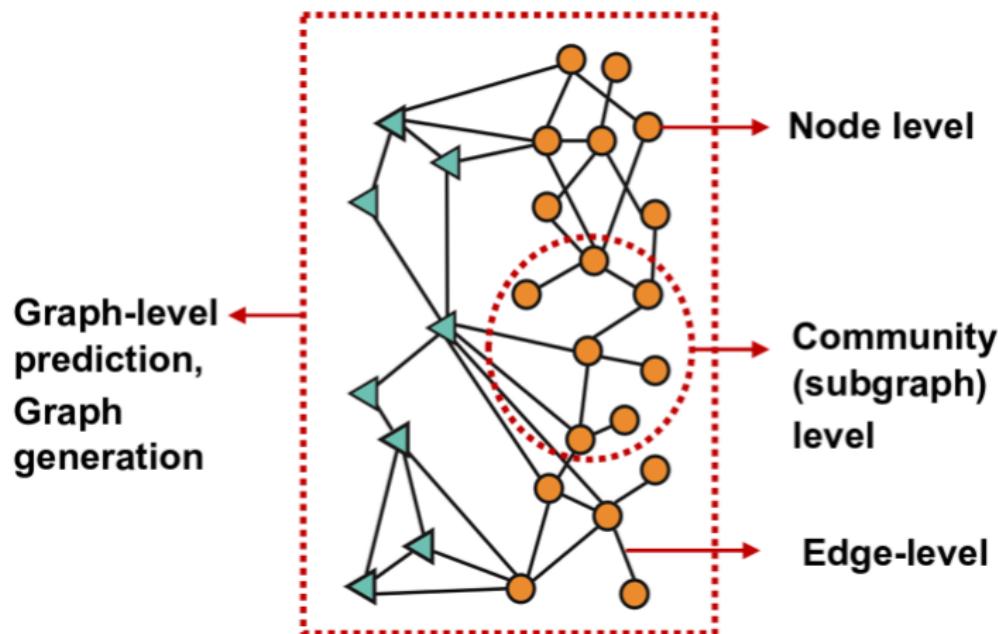
## 2 Types of Tasks

### 3 Why is Graph Deep Learning Challenging?

### 4 Deep Learning on Graphs

### 5 Classical GNN Layers

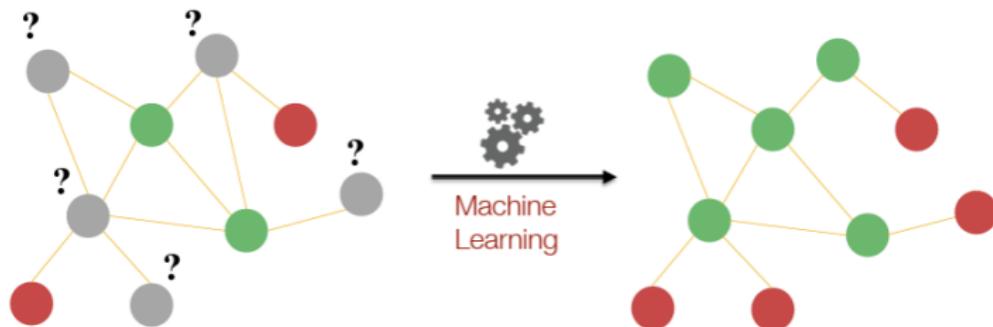
# Types of Tasks



# Node-Level Tasks

## Goal

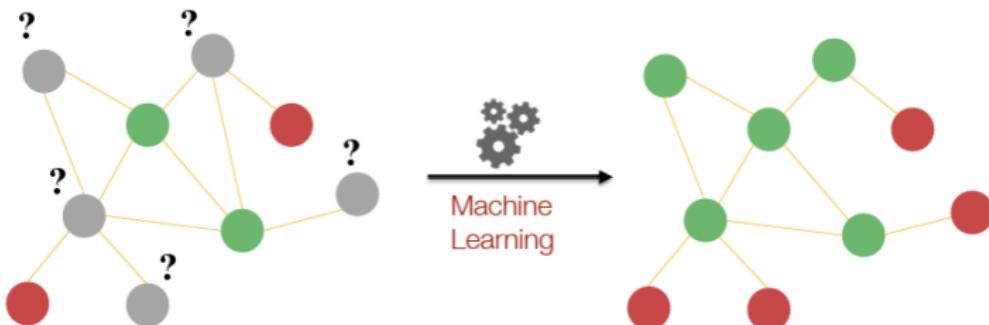
Predict a label, category, or attribute for each individual node within a graph.



# Node-Level Tasks

## Goal

Predict a label, category, or attribute for each individual node within a graph.



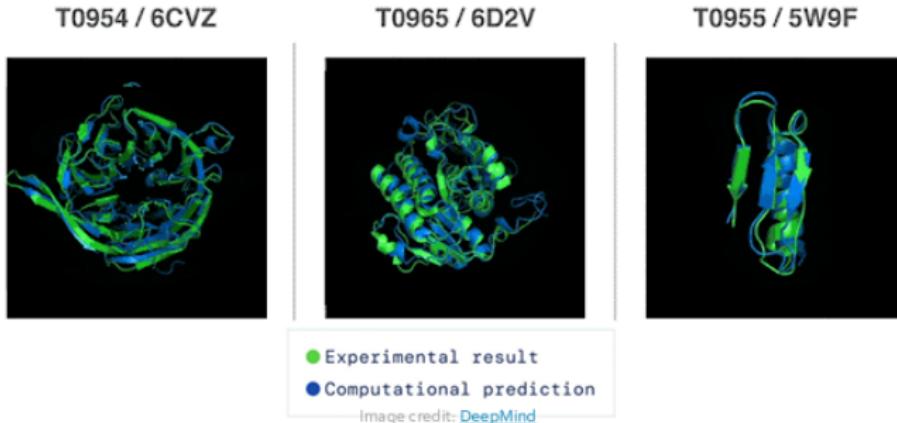
Common tasks include:

- **Node classification** (e.g., user type, document topic)
- **Node regression** (e.g., ranking score, age prediction)

# Example: Node-Level Regression for Protein Folding

## Breakthrough: Predicting Protein Structures with GNNs

- Recent advances in deep learning, notably by DeepMind's **AlphaFold**, have revolutionized the protein folding problem.
- Proteins can be modeled as **graphs**, where:
  - Nodes:** Amino acids (residues) in the protein sequence
  - Edges:** Spatial proximity or learned interaction between residues
  - Task:** For each node, **predict its 3D coordinates ( $x, y, z$ )**



# Example: Node-Level Regression for Protein Folding

## Why Graph Neural Networks?

- Proteins are naturally irregular, flexible, and structured — ideal for graphs
- Their folded structure depends on long-range, non-local interactions — not just sequence order.
- GNNs enable message passing between residues to learn spatial dependencies
- Enables accurate modeling of physical constraints and folding behavior

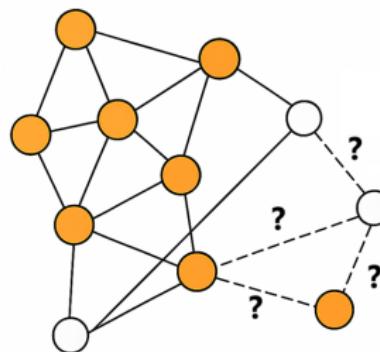
## Why It Matters

AlphaFold achieved near-experimental accuracy in predicting protein 3D structures, solving a 50-year grand challenge in biology — enabled in part by graph-based learning.

# Edge-Level Tasks

## Goal

- Predict whether an **edge exists** between two nodes — even if it is currently missing or unknown.
- At test time: Given a set of node pairs with **no existing edge**, assign a score to each pair and **rank** them.
- The top- $K$  ranked pairs are predicted as new or likely edges.
- It is essentially a **binary classification / scoring task** over node pairs.



# Example: Edge-Level Prediction for Recommender Systems

## Problem Setting

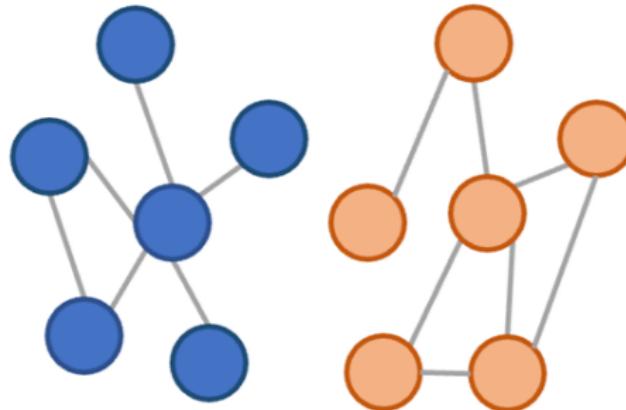
- Model user-item interactions as a bipartite graph:
  - Nodes:** Users and items (e.g., movies, books, products)
  - Edges:** Observed interactions (e.g., rating, click, purchase)
- Goal: **Predict new edges** between users and items they haven't interacted with yet.
- This is a classic **link prediction** task on a bipartite graph.



# Graph-Level Tasks

## What Are Graph-Level Tasks?

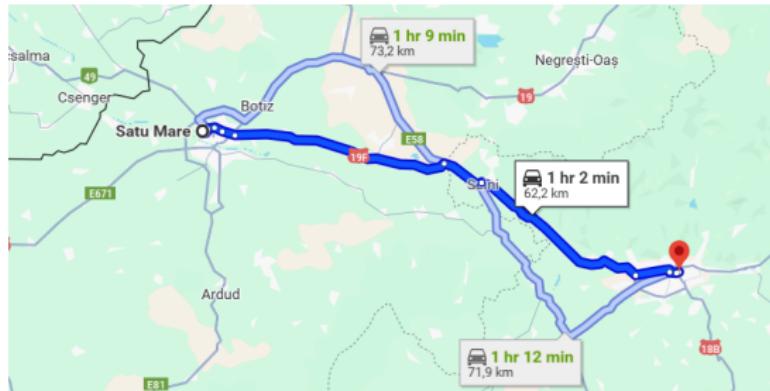
- The goal is to predict a label, property, or output for an entire graph — not individual nodes or edges.
- The model learns a **global representation (embedding)** of the entire graph structure and features.



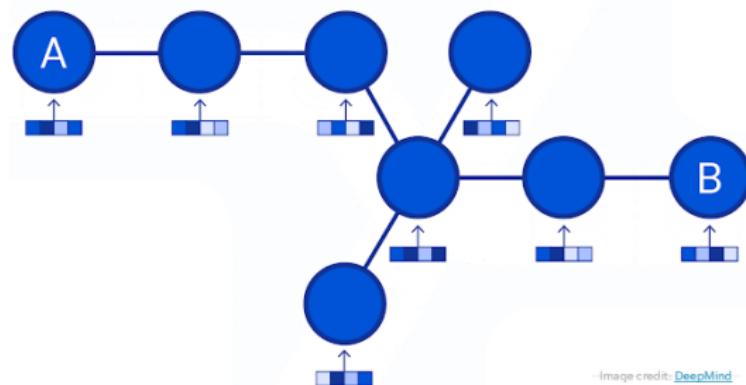
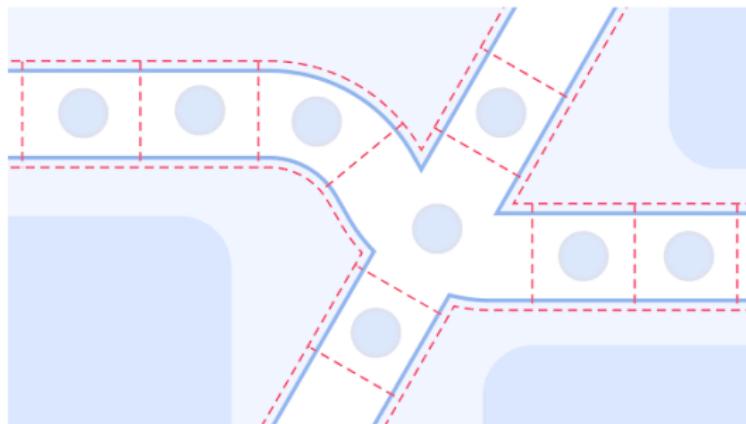
# Example: ETA Predictions in Google Maps

## Problem Setting

- Road Network as a Graph
  - ▶ **Nodes:** Represent road segments or intersections.
  - ▶ **Edges:** Connectivity between road segments / intersections (roads).
- Task: Predict the **Time of Arrival (ETA)** at the destination given the starting point.
- Use **real-time traffic data** (traffic congestion, accidents) to adjust the travel time along each road segment.



# Example: ETA Predictions in Google Maps



-Image credit: [DeepMind](#) -

1 Introduction

2 Types of Tasks

3 Why is Graph Deep Learning Challenging?

4 Deep Learning on Graphs

5 Classical GNN Layers

# Why is Graph Deep Learning Challenging?

## Irregular structure

Networks are inherently complex

### Irregular Structure of Graphs

- Unlike images or sequences, graphs have a **highly irregular structure**.
- Traditional models like CNNs or RNNs rely on fixed input shapes (e.g., images with consistent height and width, or padded sequences).

3	1	4	1	5	0	0	0	0	0
7	8	9	0	0	0	0	0	0	0
2	6	0	0	0	0	0	0	0	0

1	2	3	4
5	6	7	8
9	10	11	12
9	11	12	12



0	0	0	0	0	0	0
0	1	3	7	9	0	0
0	4	6	10	11	0	0
0	11	12	13	12	0	0
0	14	11	14	15	0	0
0	0	0	0	0	0	0

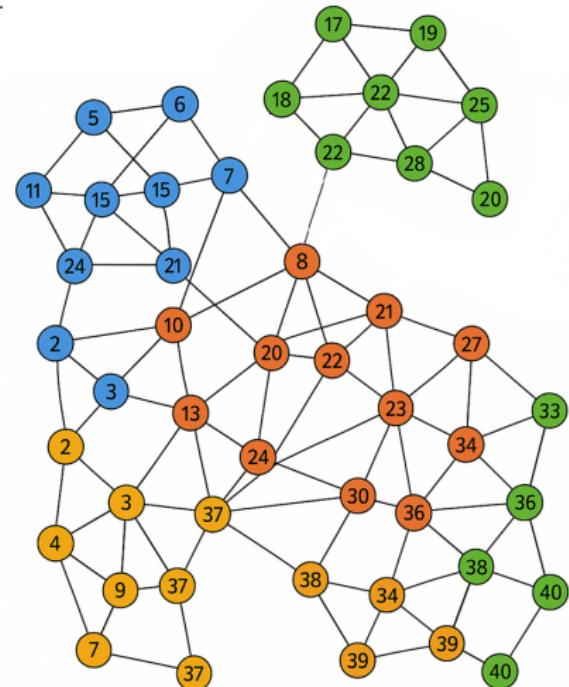
# Why is Graph Deep Learning Challenging?

## Irregular structure

- Graphs, however, can:
  - ▶ Vary widely in the number of **nodes**.
  - ▶ Have **nodes with varying degrees of connectivity**.

### Key Implications

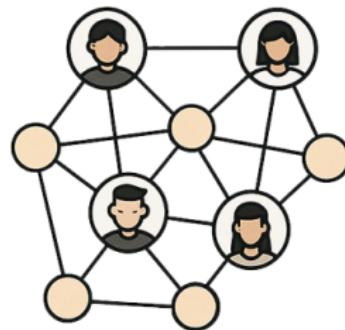
- This irregularity makes it difficult to apply standard neural operations (e.g., **convolution** or **pooling**), which assume consistent input sizes.
- Developing operations that can effectively work with these varying structures is a fundamental challenge in graph deep learning.



# Why is Graph Deep Learning Challenging?

## Interdependence of Data Points

- Traditional machine learning models typically assume that instances are **independent** of each other.
- However, in graph data, this assumption doesn't hold:
  - ▶ Each **node** (data point) is inherently **connected** to others.



### Why This Matters

- The interconnected nature of graph data requires models that can handle these **dependencies**.
- Traditional models are not designed to account for these relationships.
- We need models that can learn from the *interactions* and *dependencies* between nodes.

# Why is Graph Deep Learning Challenging?

## Permutation Invariance

**Traditional** data: fixed order

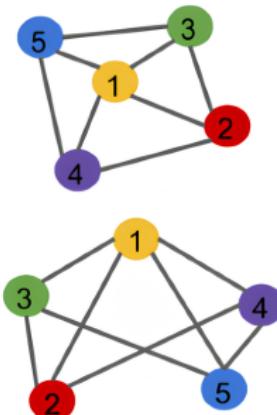
- **Images, text:** the order of data is meaningful
- **CNNs and RNNs** process data in a fixed order.

**Graph** data: **NO** fixed order

- In **graphs**, nodes (entities) have no inherent order.
- **Permutation invariance:** edges connect nodes, but the order of the nodes doesn't change the graph.

### Challenges for Graph Models

- The model should give the same output regardless of the order in which nodes and edges are presented.
- Shuffling node IDs does not change the graph structure, posing challenges for traditional models that rely on fixed input structures.



# Why is Graph Deep Learning Challenging?

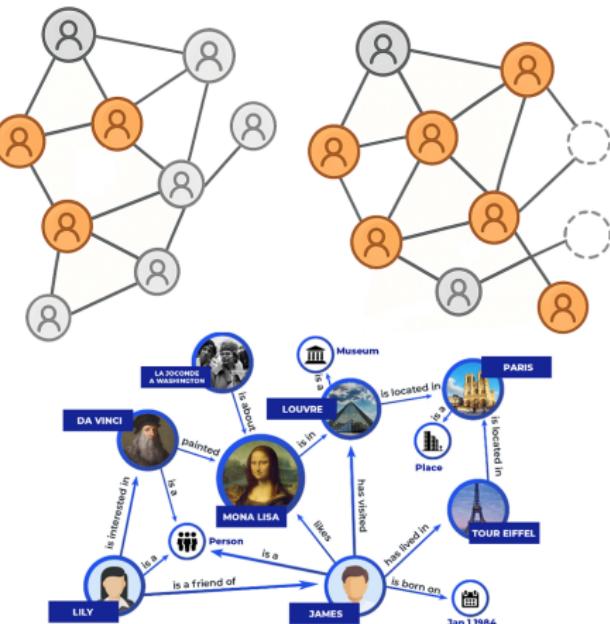
## Dynamic and Multimodal Graphs

- **Dynamic** graphs:

- ▶ Real-world graphs are evolving over time: nodes and edges can be added, removed, or updated.

- **Multimodal** graphs:

- ▶ Contain different types of nodes and edges
- ▶ Different node and edge types can have various associated features, such as:
  - **User nodes** with features like age, location, interests
  - **Transaction edges** with attributes like purchase amount, time, and product type.



## 1 Introduction

## 2 Types of Tasks

## 3 Why is Graph Deep Learning Challenging?

## 4 Deep Learning on Graphs

- Basics of Deep Learning
- Setup
- How Can We Apply Neural Networks to Graphs?
- Graph Neural Networks
- A Single GNN Layer

## 5 Classical GNN Layers

# Basics of Deep Learning

**Supervised Learning:** predict an output  $\mathbf{y}$  from an input  $\mathbf{x}$ .

- The input  $\mathbf{x}$  can be vectors, sequences, matrices, or **graphs**.
- The task is framed as an optimization problem:

$$\min_{\theta} \mathcal{L}(\mathbf{y}, f(\mathbf{x}))$$

- $\theta$ : a set of **parameters** we optimize
  - ▶ could contain one or more scalar, vector, matrices, etc.
- $\mathcal{L}$ : **Loss Function**. Examples:
  - ▶ L2 loss:

$$\mathcal{L}(\mathbf{y}, f(\mathbf{x})) := \|\mathbf{y} - f(\mathbf{x})\|^2$$

- ▶ Cross Entropy:

$$\mathcal{L}(\mathbf{y}, f(\mathbf{x})) := - \sum_{i=1}^C y_i \log f(x)_i$$

where  $\mathbf{y} \in \{0, 1\}^C$  is a categorical vector (one-hot encoding).

- ▶ Other loss function: L1 loss, max margin (hinge loss), etc.

# Basics of Deep Learning

- $f$  can be a simple linear layer, an MLP, CNN, or other neural networks (e.g., a GNN later).
  - ▶ A multi-layer perceptron (MLP) is a neural network composed of stacked layers:

$$f_i(x) = \sigma(\mathbf{W}x + \mathbf{b})$$

- $\mathbf{W}$  and  $\mathbf{b}$  are learnable parameters (weights and biases).
- $\sigma$  is the activation function (e.g., ReLU, Sigmoid).

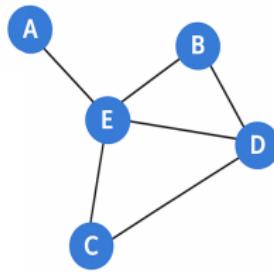
- **Optimization:** We solve the optimization problem using gradients:

$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}$$

- ▶  $\eta$  is the learning rate.
- ▶  $\nabla_{\theta} \mathcal{L} \rightarrow$  directional derivative in the direction of largest increase
- ▶ Use stochastic gradient descent (SGD) to optimize  $\mathcal{L}$  for weights  $\theta$  over many iterations.
- ▶ Other higher order optimisers: RMSprop, Adam, Adagrad, etc.

# Setup

- Assume we have a graph  $G = (V, E)$ , where:
  - $V \rightarrow$  the set of graph nodes  $V = \{v_1, \dots, v_N\}$ ,  $N = \text{num. nodes}$
  - $E \rightarrow$  the edge set, each edge defined as  $(v_i, v_j)$
- $X = [x_1, \dots, x_N] \rightarrow$  node feature set,  $X \in \mathbb{R}^{NxD}$ ,  $D = \text{dimension of } x_i$
- $A \in \mathbb{R}^{N \times N} \rightarrow$  adjacency matrix:
  - binary: 1's and 0's
  - weighted: each edge has a strength  $e_{ij} \in R$ :
    - $e_{ij} > 0$ , if  $(v_i, v_j) \in E$
    - $e_{ij} = 0$ , if  $(v_i, v_j) \notin E$
- $D_{ii} = \sum_j A_{ij}$  degree matrix



	A	B	C	D	E
A	0	0	0	0	1
B	0	0	0	1	1
C	0	0	0	1	1
D	0	1	1	0	0
E	1	1	1	1	0

Adjacency matrix A

	A	B	C	D	E
A	1	0	0	0	0
B	0	2	0	0	0
C	0	0	2	0	0
D	0	0	0	3	0
E	0	0	0	0	4

Degree matrix D

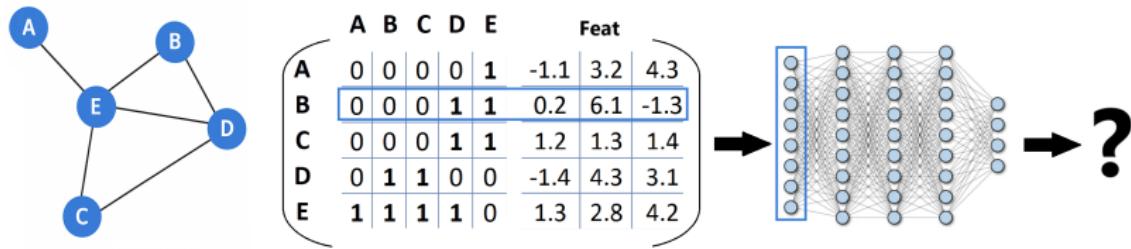
	A	B	C
A	-1.1	3.2	4.3
B	0.2	6.1	-1.3
C	1.2	1.3	1.4
D	-1.4	4.3	3.1
E	1.3	2.8	4.2

Feature vector X

# How Can We Apply Neural Networks to Graphs?

## Naive Approach for Graph Representation

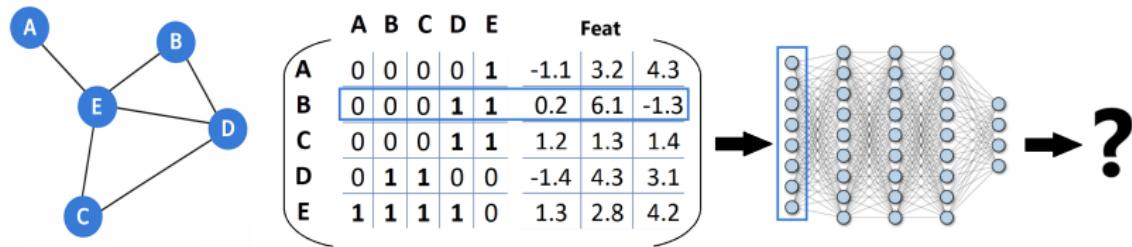
- Join the *adjacency matrix* and *node features*
- Feed them into a deep neural network.



# How Can We Apply Neural Networks to Graphs?

## Naive Approach for Graph Representation

- Join the *adjacency matrix* and *node features*
- Feed them into a deep neural network.

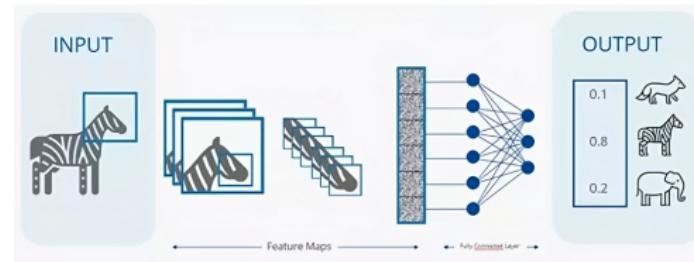
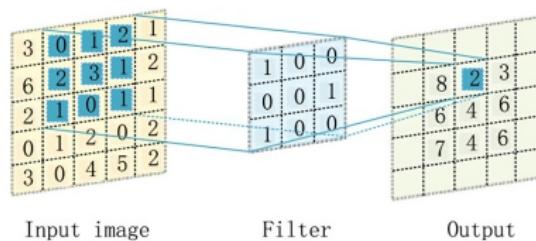


### Problems with This Approach

- $O(|V|)$  parameters
- Not applicable to graphs of different sizes.
- Sensitive to node ordering.

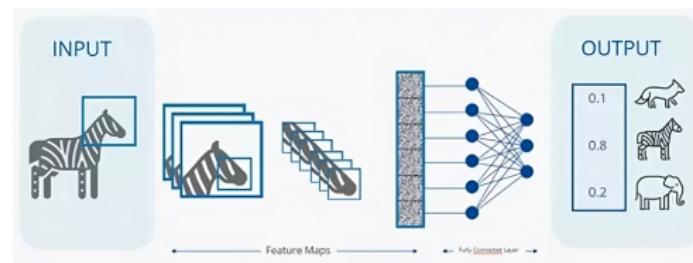
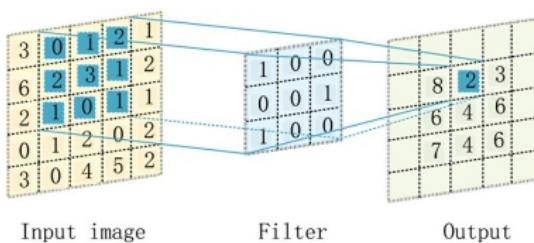
# How Can We Apply Neural Networks to Graphs?

## Idea: Convolutional Networks



# How Can We Apply Neural Networks to Graphs?

## Idea: Convolutional Networks



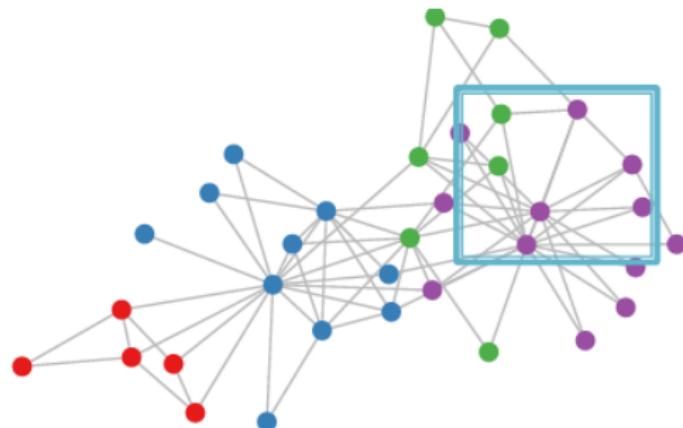
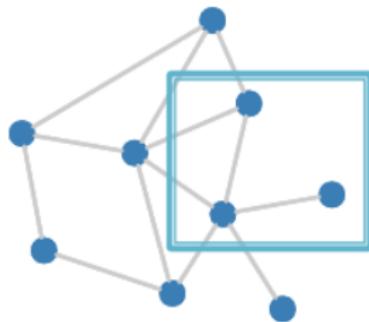
## Goals

- Generalize convolutions beyond simple lattices
- Leverage node features/attributes.

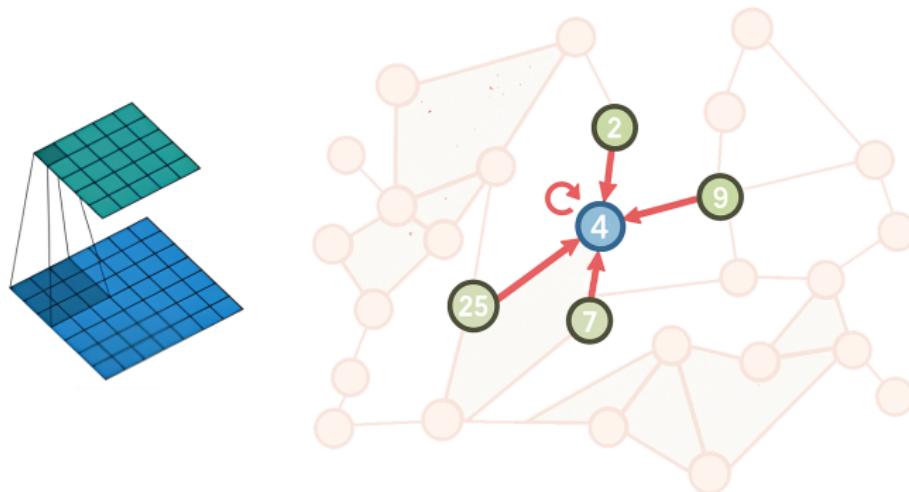
# From Images to Graphs

## What if Our Data Looks Like This?

- *Graph-structured data* lacks a fixed notion of locality.
- There's no *sliding window* over graphs.
- Graphs are **permutation-invariant**.



# From Images to Graphs



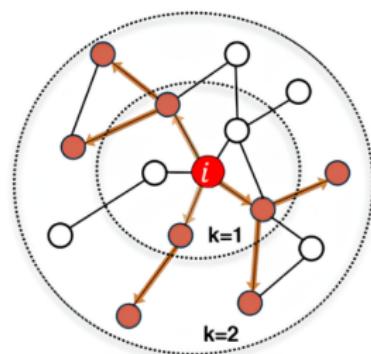
## Idea

- **Transform** neighbor information and **combine** it:
  - ▶ Transform *messages*  $h_i$  from neighbors:  $W_i h_i$
  - ▶ Sum them up:  $\sum_i W_i h_i$

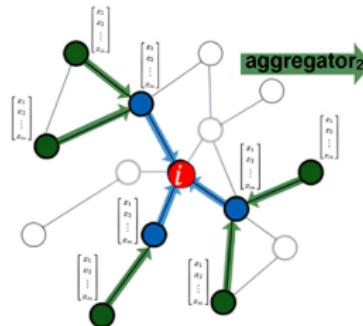
# Graph Neural Networks

## Key Ideas

- The neighborhood of a node defines a **computation graph**
- Learn how to **propagate** information across the graph to compute node features



Determine node computation graph

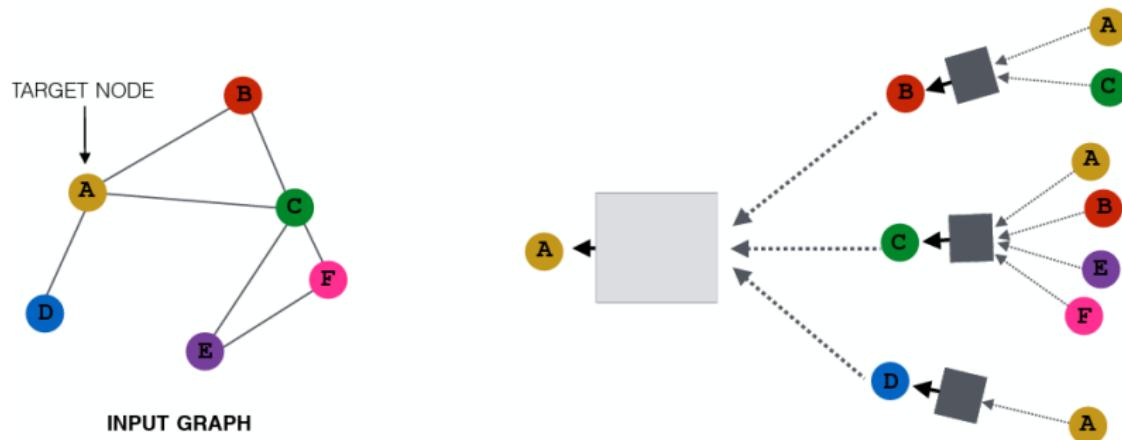


Propagate and transform information

# Idea: Aggregating Information from Neighbors

## Key Idea

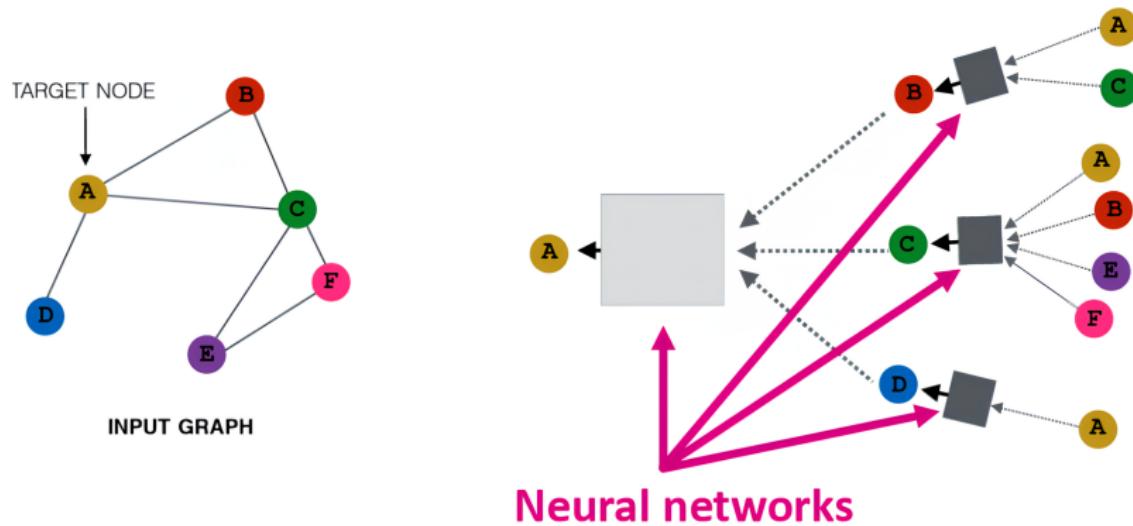
- Generate node embeddings based on **local network neighborhoods**



# Idea: Aggregating Information from Neighbors

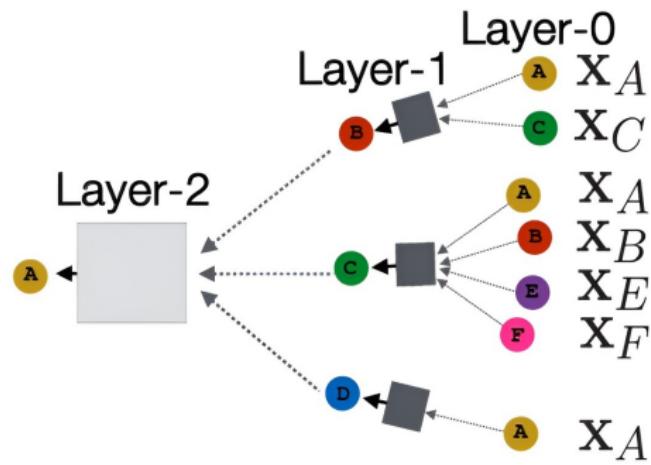
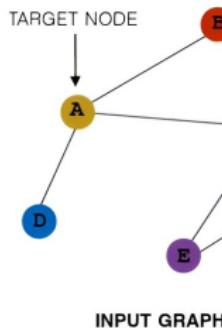
## Key Idea

- Generate node embeddings based on **local network neighborhoods**



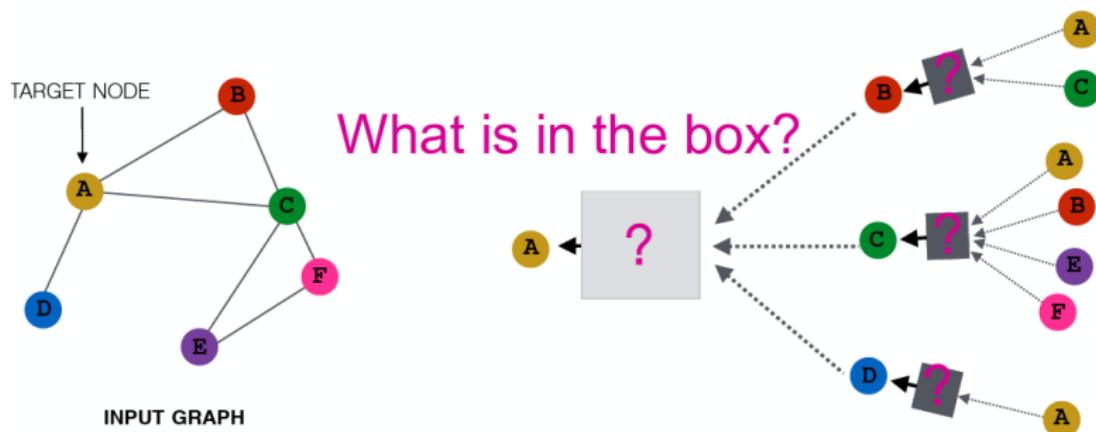
# Network Depth

- The network can have an arbitrary depth.
- At each layer, nodes have **embeddings** (features).
- Layer 0:** embedding of node  $v$  is its input feature  $x_v$ .
- Layer-k:** embeddings  $h_v$  contain information from nodes that are  $K$  hops away in the graph.



# Neighborhood Aggregation

**Key distinction between GNNs:** How different approaches aggregate information across layers.



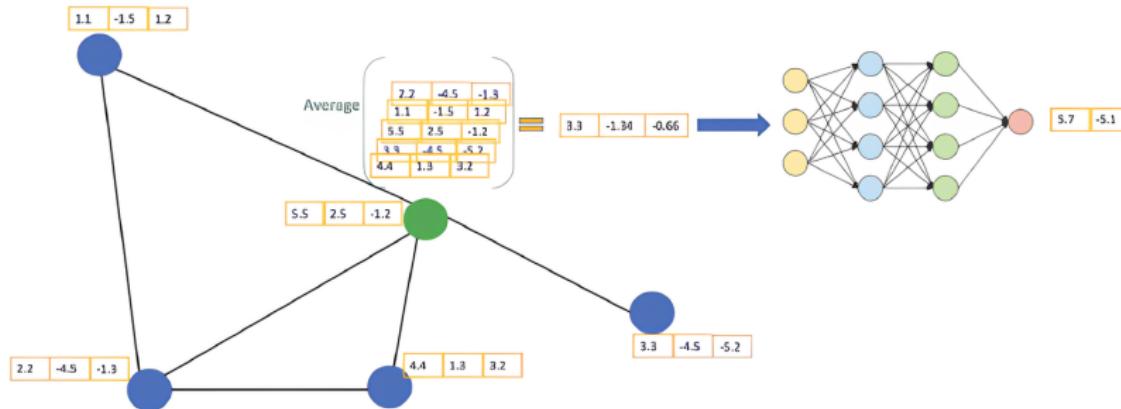
# Neighborhood Aggregation

**Basic approach?**

# Neighborhood Aggregation

## Basic approach?

Average information from neighbors and apply a neural network.



# Graph Deep Encoder: The Math

## Basic approach?

Average information from neighbors and apply a **neural network**.

$$h_v^{(0)} = x_v$$

Initial 0-th layer embeddings  
are equal to node features

# Graph Deep Encoder: The Math

## Basic approach?

Average information from neighbors and apply a neural network.

$$h_v^{(0)} = x_v$$

Initial 0-th layer embeddings  
are equal to node features



$$h_v^{(k+1)} = \sigma \left( W_k^{(0)} h_v^{(k)} + W_k^{(1)} \sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|} \right), \forall k \in \{0, \dots, K-1\}$$

# Graph Deep Encoder: The Math

## Basic approach?

Average information from neighbors and apply a **neural network**.

$$h_v^{(0)} = x_v$$

Initial 0-th layer embeddings  
are equal to node features

$$h_v^{(k+1)} = \sigma \left( W_k^{(0)} h_v^{(k)} + W_k^{(1)} \sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|} \right), \forall k \in \{0, \dots, K-1\}$$

Embedding of  $v$  at layer  $k$

# Graph Deep Encoder: The Math

## Basic approach?

**Average** information from neighbors and apply a **neural network**.

$$h_v^{(0)} = x_v$$

Initial 0-th layer embeddings  
are equal to node features

Embedding of  $v$  at layer  $k$

$$h_v^{(k+1)} = \sigma \left( W_k^{(0)} h_v^{(k)} + W_k^{(1)} \sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|} \right), \forall k \in \{0, \dots, K-1\}$$

Average of neighbor's previous  
layer embeddings

# Graph Deep Encoder: The Math

## Basic approach?

Average information from neighbors and apply a **neural network**.

$$h_v^{(0)} = x_v$$

Initial 0-th layer embeddings  
are equal to node features

Non-linearity (e.g., ReLU)

$$h_v^{(k+1)} = \sigma \left( W_k^{(0)} h_v^{(k)} + W_k^{(1)} \sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|} \right), \forall k \in \{0, \dots, K-1\}$$

Embedding of  $v$  at layer  $k$

Average of neighbor's previous  
layer embeddings

# Graph Deep Encoder: The Math

## Basic approach?

Average information from neighbors and apply a **neural network**.

$$h_v^{(0)} = x_v$$

Initial 0-th layer embeddings  
are equal to node features

Non-linearity (e.g., ReLU)

$$h_v^{(k+1)} = \sigma \left( W_k^{(0)} h_v^{(k)} + W_k^{(1)} \sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|} \right), \forall k \in \{0, \dots, K-1\}$$

Embedding of  $v$  at layer  $k$

Average of neighbor's previous  
layer embeddings

Number of layers

# Graph Deep Encoder: The Math

## Basic approach?

Average information from neighbors and apply a **neural network**.

$$h_v^{(0)} = x_v$$

Initial 0-th layer embeddings  
are equal to node features

Non-linearity (e.g., ReLU)

$$h_v^{(k+1)} = \sigma \left( W_k^{(0)} h_v^{(k)} + W_k^{(1)} \sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|} \right), \forall k \in \{0, \dots, K-1\}$$

Average of neighbor's previous  
layer embeddings

Embedding of  $v$  at layer  $k$

Number of layers

$$z_v = h_v^{(K)}$$

Embedding after  $K$  layers of  
neighborhood aggregation

# Efficient Aggregation Using Matrix Operations

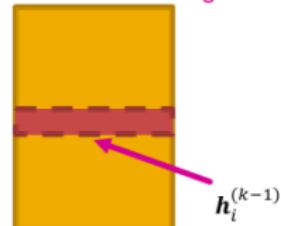
**Many aggregations can be performed efficiently by (sparse) matrix operations**

$$\text{Let } H^{(k)} = [h_1^{(k)} \quad h_2^{(k)} \quad \dots \quad h_{|V|}^{(k)}]^T$$

Then the node feature aggregation can be written as:

$$\sum_{v \in N(u)} h_v^{(k)} = A_{:,v} H^{(k)}$$

Matrix of hidden embeddings  $H^{(k-1)}$



Let  $D$  be the diagonal degree matrix:

$$D_{v,v} = \text{Deg}(v) = |N(v)|$$

Note: The inverse of  $D$  is also diagonal.

# Efficient Aggregation Using Matrix Operations

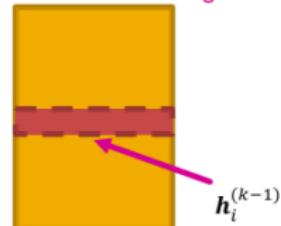
**Many aggregations can be performed efficiently by (sparse) matrix operations**

$$\text{Let } H^{(k)} = [h_1^{(k)} \quad h_2^{(k)} \quad \dots \quad h_{|V|}^{(k)}]^T$$

Then the node feature aggregation can be written as:

$$\sum_{v \in N(u)} h_v^{(k)} = A_{:,v} H^{(k)}$$

Matrix of hidden embeddings  $H^{(k-1)}$



Let  $D$  be the diagonal degree matrix:

$$D_{v,v} = \text{Deg}(v) = |N(v)|$$

Note: The inverse of  $D$  is also diagonal.

Therefore:

$$\sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|} \implies \mathbf{H}^{(k+1)} = \mathbf{D}^{-1} \mathbf{A} \mathbf{H}^{(k)}$$

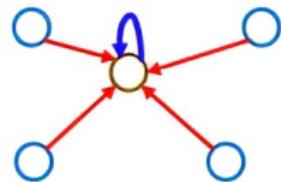
# GNN Update Rule: Matrix Form

## Update Equation

$$H^{(k+1)} = \sigma \left( \tilde{A} H^{(k)} W_k^{(1)^T} + H^{(k)} W_k^{(0)^T} \right)$$

where

- $\tilde{A} = D^{-1}A$ : normalized adjacency matrix
- $H^{(k)} = [h_1^{(k)} \dots h_{|V|}^{(k)}]^\top$ : node features at layer  $k$



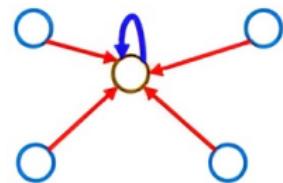
# GNN Update Rule: Matrix Form

## Update Equation

$$H^{(k+1)} = \sigma \left( \tilde{A} H^{(k)} W_k^{(1)^T} + H^{(k)} W_k^{(0)^T} \right)$$

where

- $\tilde{A} = D^{-1}A$ : normalized adjacency matrix
- $H^{(k)} = [h_1^{(k)} \dots h_{|V|}^{(k)}]^\top$ : node features at layer  $k$



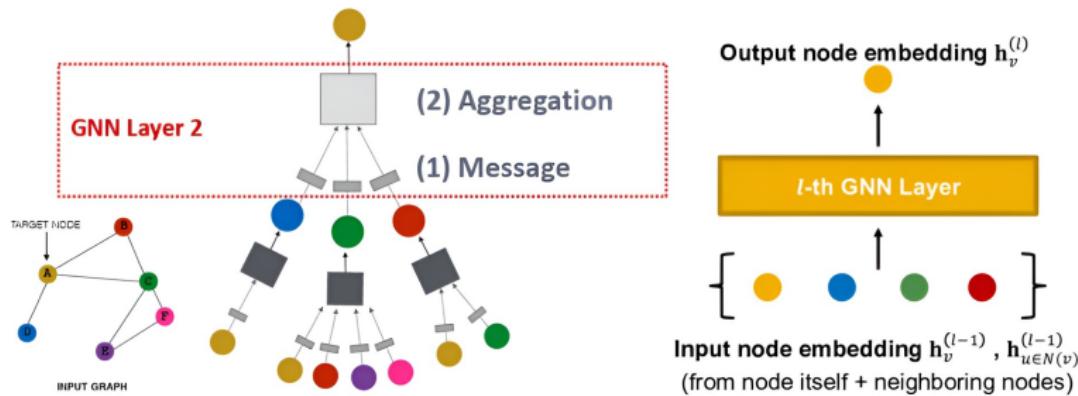
## Efficiency Note

$A$  is sparse  $\Rightarrow$  enables efficient sparse matrix multiplication

# A Single GNN Layer

## Key Idea

- Compress a set of vectors into a single vector
- **GNN Layer = Message + Aggregation**



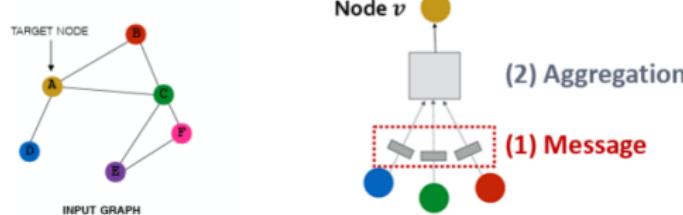
# A Single GNN Layer: Message Computation

## Message Computation

$$\mathbf{m}_u^{(k)} = \text{MSG}^{(k)}(\mathbf{h}_u^{(k-1)})$$

- **Intuition:** Each node creates a message to be sent to its neighbors.
- **Example:** Linear transformation

$$\mathbf{m}_u^{(k)} = \mathbf{W}_1^{(k)} \mathbf{h}_u^{(k-1)}$$



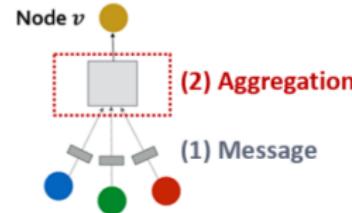
# A Single GNN Layer: Aggregation

## Aggregation

$$\mathbf{h}_v^{(k)} = \text{AGG}^{(k)} \left( \left\{ \mathbf{m}_u^{(k)} \mid u \in \mathcal{N}(v) \right\} \right)$$

- **Intuition:** Node  $v$  aggregates messages from its neighbors  $u \in \mathcal{N}(v)$
- **Examples of AGG:** `sum()`, `mean()`, or `max()`
- **Example:**

$$\mathbf{h}_v^{(k)} = \sum_{u \in \mathcal{N}(v)} \mathbf{m}_u^{(k)}$$



# A Single GNN Layer: Aggregation Issue and Solution

## Issue: Node's Own Information May Be Lost

- Aggregation only over neighbors:  $\mathbf{h}_v^{(k)} = \text{AGG}^{(k)} \left( \left\{ \mathbf{m}_u^{(k)} \mid u \in \mathcal{N}(v) \right\} \right)$
- **Problem:**  $\mathbf{h}_v^{(k)}$  does not directly depend on node  $v$ 's own features.

# A Single GNN Layer: Aggregation Issue and Solution

## Issue: Node's Own Information May Be Lost

- Aggregation only over neighbors:  $\mathbf{h}_v^{(k)} = \text{AGG}^{(k)} \left( \left\{ \mathbf{m}_u^{(k)} \mid u \in \mathcal{N}(v) \right\} \right)$
- Problem:**  $\mathbf{h}_v^{(k)}$  does not directly depend on node  $v$ 's own features.

## Solution: Include Self-Messages

- (1) Message from node  $v$  itself:**  $\mathbf{m}_v^{(k)} = \mathbf{W}_0^{(k)} \mathbf{h}_v^{(k-1)}$   
(separate from  $\mathbf{W}_1^{(k)}$  used for neighbors)
- (2) Aggregation:**
  - First aggregate from neighbors:

$$\mathbf{h}_{\mathcal{N}(v)}^{(k)} = \text{AGG}^{(k)} \left( \left\{ \mathbf{m}_u^{(k)} \mid u \in \mathcal{N}(v) \right\} \right)$$

- Then combine with self-message:

$$\mathbf{h}_v^{(k)} = \text{CONCAT} \left( \mathbf{h}_{\mathcal{N}(v)}^{(k)}, \mathbf{m}_v^{(k)} \right)$$

- 1 Introduction
- 2 Types of Tasks
- 3 Why is Graph Deep Learning Challenging?
- 4 Deep Learning on Graphs
- 5 Classical GNN Layers
  - Graph Convolutional Networks
  - GraphSAGE

# Classical GNN Layers: GCN

## Graph Convolutional Network (GCN) Update Rule

$$\mathbf{h}_v^{(k)} = \sigma \left( \sum_{u \in \mathcal{N}(v) \cup \{v\}} \frac{1}{\sqrt{d_u d_v}} \mathbf{W}^{(k)} \mathbf{h}_u^{(k-1)} \right)$$

- **Message:**

$$\mathbf{m}_u^{(k)} = \mathbf{W}^{(k)} \mathbf{h}_u^{(k-1)}$$

Each neighbor (and the node itself) sends a message using shared weight matrix  $\mathbf{W}^{(k)}$ .

- **Aggregation:**

- ▶ Normalize messages by node degrees  $d_u, d_v$
- ▶ Sum messages from neighbors **and the node itself**
- ▶ Apply non-linearity  $\sigma$  (e.g., ReLU)

- **Note:** The input graph is assumed to include self-loops.

This formulation is from Kipf & Welling (2017).

# Classical GNN Layers: GraphSAGE

## GraphSAGE Update Rule

$$\mathbf{h}_v^{(k)} = \sigma \left( \mathbf{W}^{(k)} \cdot \text{CONCAT} \left( \mathbf{h}_v^{(k-1)}, \text{AGG}^{(k)} \left( \left\{ \mathbf{h}_u^{(k-1)} \mid u \in \mathcal{N}(v) \right\} \right) \right) \right)$$

## Two-Stage Aggregation

- **Stage 1: Neighborhood Aggregation**

$$\mathbf{h}_{\mathcal{N}(v)}^{(k)} = \text{AGG}^{(k)} \left( \left\{ \mathbf{h}_u^{(k-1)} \mid u \in \mathcal{N}(v) \right\} \right)$$

- **Stage 2: Combine with Node Itself**

$$\mathbf{h}_v^{(k)} = \sigma \left( \mathbf{W}^{(k)} \cdot \text{CONCAT} \left( \mathbf{h}_v^{(k-1)}, \mathbf{h}_{\mathcal{N}(v)}^{(k)} \right) \right)$$

This formulation is from Hamilton et al., (2017).

It formulation supports inductive learning on unseen graphs.

# Classical GNN Layers: GraphSAGE Neighbor Aggregation

- **Mean:** Take an average of neighbor features

$$\text{AGG}^{(k)} = \underbrace{\sum_{u \in \mathcal{N}(v)} \frac{\mathbf{h}_u^{(k-1)}}{|\mathcal{N}(v)|}}_{\text{Aggregation}} \underbrace{\quad}_{\text{Message Computation (Identity)}}$$

- **Pool:** Apply transformation before aggregation (e.g., MLP)

$$\text{AGG}^{(k)} = \underbrace{\text{Mean} \left( \left\{ \underbrace{\text{MLP}^{(k)}(\mathbf{h}_u^{(k-1)})}_{\text{Message Computation}} \mid u \in \mathcal{N}(v) \right\} \right)}_{\text{Aggregation}}$$

- **LSTM:** Order-sensitive aggregation over neighbors

$$\text{AGG}^{(k)} = \underbrace{\text{LSTM}^{(k)} \left( \left[ \mathbf{h}_u^{(k-1)} \mid u \in \pi(\mathcal{N}(v)) \right] \right)}_{\text{Aggregation (learned over sequence)}}$$

# Questions?

---